

Definition of a problem statement and a short outline of the implementation

Problem statement

Description of data acquisition / how it was collected (by you or the publisher of the data)

1. Data preparation (general)
2. Natural Language Processing
3. Supervised / Unsupervised ML
4. Network Analysis

# M2 - Exam - Network and NLP

Code ▼

Malika Kuhlman & Michael Dybdahl

10/24/2019

**OBS** The notebook cannot be run in google Colab because of the *topicmodels* package.

Loading the *stop\_words* that we will use in the tokenization.

Hide

```
set.seed(9)
data("stop_words")
```

## Definition of a problem statement and a short outline of the implementation

The link to our Github: [https://github.com/michaeldybdahl/M2\\_exam](https://github.com/michaeldybdahl/M2_exam)  
([https://github.com/michaeldybdahl/M2\\_exam](https://github.com/michaeldybdahl/M2_exam))

## Problem statement

How are the words in the text-reviews and based on a tidy format of the reviews, can we then compute any topics based on this? Can we compute a model that predict if a review is rated as a good review, based or linked to the NLP results?

## Description of data acquisition / how it was collected (by you or the publisher of the data)

The dataset contains reviews of some of amazons products, the dataset can be reached on this website <https://data.world/datafiniti/consumer-reviews-of-amazon-products>  
(<https://data.world/datafiniti/consumer-reviews-of-amazon-products>)

There are several datasets about the same topic, but with different amount of observations, we chose the dataset with 5000 observations. The dataset contains 24 variables and 5.000 observations (reviews).

Hide

```
reviews <- read_csv("https://query.data.world/s/foqg5o75hazenbwqdug534atoqiy
p3")
```

Have a look at the data structure

Hide

```
glimpse(reviews)
```

```
## Observations: 5,000
## Variables: 24
## $ id <chr> "AVqVGZNvQmLgsOJE6eUY", "AVqVGZNvQmLgsOJE6eU...
## $ dateAdded <dtm> 2017-03-03 16:56:05, 2017-03-03 16:56:05, 2...
## $ dateUpdated <dtm> 2018-10-25 16:36:31, 2018-10-25 16:36:31, 2...
## $ name <chr> "Amazon Kindle E-Reader 6\" Wifi (8th Genera...
## $ asins <chr> "B00ZV9PXP2", "B00ZV9PXP2", "B00ZV9PXP2", "B...
## $ brand <chr> "Amazon", "Amazon", "Amazon", "Amazon", "Ama...
## $ categories <chr> "Computers,Electronics Features,Tablets,Elec...
## $ primaryCategories <chr> "Electronics", "Electronics", "Electronics",...
## $ imageURLs <chr> "https://piscs.bbystatic.com/image2/BestBuy...
## $ keys <chr> "allnewkindleereaderblack6glarefreetouchscre...
## $ manufacturer <chr> "Amazon", "Amazon", "Amazon", "Amazon", "Ama...
## $ manufacturerNumber <chr> "B00ZV9PXP2", "B00ZV9PXP2", "B00ZV9PXP2", "B...
## $ reviews.date <dtm> 2017-09-03 00:00:00, 2017-06-06 00:00:00, 2...
## $ reviews.dateAdded <lgl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ reviews.dateSeen <chr> "2018-05-27T00:00:00Z,2017-09-18T00:00:00Z,2...
## $ reviews.doRecommend <lgl> FALSE, TRUE, TRUE, TRUE, TRUE, FALSE, TRUE, ...
## $ reviews.id <dbl> NA, NA, NA, 177283626, NA, NA, 187043823, NA...
## $ reviews.numHelpful <dbl> 0, 0, 0, 3, 0, 0, 0, 0, 0, 3, 0, 0, 0, 0,...
## $ reviews.rating <dbl> 3, 5, 4, 5, 5, 5, 5, 4, 5, 5, 5, 5, 5, 5,...
## $ reviews.sourceURLs <chr> "http://reviews.bestbuy.com/3545/5442403/rev...
## $ reviews.text <chr> "I thought it would be as big as small paper...
## $ reviews.title <chr> "Too small", "Great light reader. Easy to us...
## $ reviews.username <chr> "llyyue", "Charmi", "johnnyjojojo", "Kdperry...
## $ sourceURLs <chr> "https://www.newegg.com/Product/Product.aspx..."
```

# 1. Data preparation (general)

## 1.1. Data cleaning

We see that the data contains of 5000 observations and 24 variables, we are not going to use all of variables, therefore we now select the one we will use. Therefore we will only describe the variables we will use.

Hide

```
reviews <- reviews %>%
  select(reviews.username,
         dateAdded,
         name,
         primaryCategories,
         reviews.numHelpful,
         reviews.rating,
         reviews.text)
```

- **reviews.username:** contain the reviewsers username.
- **dateAdded:** coontain the date the reviews was written.
- **name:** contain the name of the product.
- **primaryCategories:** contain the PRimary category of the product, there are four different primary categories in our data; *Electronics, Electronics, Hardware, Electronics, Mediaand Office Supplies, Electronics*.
- **reviews.numHelpful:** contain a the value wether or not a reader found the review helpful.
- **reviews.text:** contain the reviews of the product.

Now we check for *NA* and drop them if any exists.

Hide

```
sum(is.na(reviews)) # First we check for NA values in the whole datasat.
```

```
## [1] 0
```

Hide

```
reviews = reviews %>%
  drop_na()
```

## 1.2. Recoding

We saw earlier that some of the variable are named *reviews.X*, this we want to change.

Hide

```
reviews <- reviews %>%
  rename(helpful = reviews.numHelpful,
         rating = reviews.rating,
         text = reviews.text,
         username = reviews.username,
         date = dateAdded)
```

Have a look at the dataset now.

Hide

```
glimpse(reviews)
```

```
## Observations: 5,000
## Variables: 7
## $ username      <chr> "llyyue", "Charmi", "johnnyjojojo", "Kdperry",...
## $ date           <dtm> 2017-03-03 16:56:05, 2017-03-03 16:56:05, 201...
## $ name           <chr> "Amazon Kindle E-Reader 6\" Wifi (8th Generati...
## $ primaryCategories <chr> "Electronics", "Electronics", "Electronics", "...
## $ helpful        <dbl> 0, 0, 0, 3, 0, 0, 0, 0, 0, 0, 3, 0, 0, 0, 0, 0...
## $ rating          <dbl> 3, 5, 4, 5, 5, 5, 5, 4, 5, 5, 5, 5, 5, 5, 3...
## $ text            <chr> "I thought it would be as big as small paper b...
```

Now we have a dataset with 4987 observations and 9 variables

For later use, we add a new *ID* number to each username, since some users have reviewed more than one product.

Hide

```
username_df <- reviews %>%
  select(username) %>%
  distinct(username,
    .keep_all = TRUE)

username_df$ID <- seq.int(nrow(username_df))
```

Now we join the new ID with the dataset.

Hide

```
reviews <- reviews %>%
  left_join(username_df, by = "username") %>%
  select(ID, everything())

reviews$ID2 <- seq.int(nrow(reviews))
```

We change the rating from *double* to *factor*.

Hide

```
reviews$rating <- as.factor(reviews$rating)
```

We change the date from *dtm* to *date*.

Hide

```
reviews$date <- as.Date(as.POSIXct(reviews$date))
```

We create a new variable called *year*, where we put the year of every review.

Hide

```
reviews <- reviews %>%
  mutate(year = as.factor(year(date)))
```

Now we want to look into which years the dataset has reviews from and how many reviews there are written in each year.

Hide

```
reviews %>%
  group_by(year) %>%
  count()
```

year	n
<fctr>	<int>
2015	44
2016	1026
2017	2495
2018	1435
4 rows	

We see we have four years

- **2015** with 44 reviews
- **2016** with 1026 reviews
- **2017** with 2495 reviews
- **2018** with 1435 reviews

## 2. Natural Language Processing

We now do all the NLP here, and afterwards we have a section containing all the Network analysis.

### 2.1. NLP - preparation

#### 2.1.1. Tokenization

In this section we are going to do a tokenization by removing *stop words*, *meaningless words*, *non-alphanumeric characters*, and so on.

First we separate every word in each review text.

Hide

```
reviews_tidy <- reviews %>%
  select(ID2, ID, text) %>%
  unnest_tokens(output = word, input = text)
```

Hide

```
head(reviews_tidy)
```

ID2	ID	word
<int>	<int>	<chr>
1	1	i
1	1	thought
1	1	it

ID2 <int>	ID word <int> <chr>
1	1 would
1	1 be
1	1 as

6 rows

Then we remove *stop words*, using the `stop_words` we loaded in the beginning.

[Hide](#)

```
reviews_tidy %<% # I use "%<%" which both assigns and pipe - instead of first assigning and then piping.
  anti_join(stop_words, by = "word") # I use anti_join instead of filter, because it should be a bit faster.
```

Now we have a look at the most used words.

[Hide](#)

```
reviews_tidy %>%
  count(word, sort = TRUE) %>% # Count "word". "sort = TRUE" means that it will sort the words in descending order of number words.
  head(20)
```

word <chr>	n <int>
tablet	1309
love	1090
easy	822
bought	785
kindle	764
amazon	694
echo	693
alexa	513
loves	506
screen	500

1-10 of 20 rows

Previous **1** 2 Next

We see that words like *tablet*, *love* and *easy* are the most used words. Since the reviews are about *Electronics* we think that's why the word *Tablet* are the most used word together with *kindle*, *alexa* and *echo* which all are names of some of the products.

After a closer look into the words, we found some words that could be removed.

[Hide](#)

```
own_stopwords <- tibble(word = c("bought", "yrs", "yokod", "yo", "xm", "wouldnt", "wouldn", "woo", "withunderstandably", "wished", "wi", "wellthis", "wellreasonably", "ample", "amazonso", "amazingly", "and"),
  lexicon = "OWN")
```

Hide

```
reviews_tidy %<>% # I use "%<>%" which both assigns and pipe - instead of first assigning and then piping.
  anti_join(own_stopwords, by = "word") # I use anti_join instead of filter, because it should be a bit faster.
```

Now we look at the top words again.

Hide

```
reviews_tidy %>%
  count(word, sort = TRUE) %>% # Count "word". "sort = TRUE" means that it will sort the words in descending order of number words.
  head(20)
```

word <chr>	n <int>
tablet	1309
love	1090
easy	822
kindle	764
amazon	694
echo	693
alexa	513
loves	506
screen	500
price	477
1-10 of 20 rows	Previous 1 2 Next

We see that the top words are different, since we removed some of our own stop words.

Here we remove numbers.

Hide

```
reviews_tidy %<>%
  mutate(word = trimws(gsub("[^\\s]*[0-9][^\\s]*", "", word, perl = T))) %>%
  filter(str_length(word) > 1) # this filter out the words that are blank.
```

Here we remove non-alphanumeric characters.

Hide

```
reviews_tidy %<%
  mutate(word = word %>% str_remove_all("[^[:alnum:]]") ) %>% # alnum = Alph
  anumeric characters.
  filter(str_length(word) > 1) # filter out words with 1 character.
```

Now we look at the top 20 of the least used words.

[Hide](#)

```
reviews_tidy %>%
  count(word, sort = TRUE) %>% # Count "word". "sort = TRUE" means that it w
  ill sort the words in descending order of number words.
  top_n(-20)
```

word <chr>	n <int>
abke	1
absent	1
absorbs	1
accelerometer	1
accept	1
acceptable	1
accessible	1
accesses	1
accessible	1
accident	1

1-10 of 2,157 rows      Previous   **1**   2   3   4   5   6   ...   216   Next

We see that the list contains words, that are only used once, this could indicate that some misspelling have happend or its not a very usefull word. Therefore we remove words used less than two times.

[Hide](#)

```
reviews_tidy = reviews_tidy %>%
  add_count(ID, word, name = "nword") %>%
  filter(nword > 1) %>%
  select(-nword)
```

Have a look again

[Hide](#)

```
reviews_tidy %>%
  count(word, sort = TRUE) %>% # Count "word". "sort = TRUE" means that it w
  ill sort the words in descending order of number words.
  top_n(-20)
```



<b>word</b>	<b>n</b>
<chr>	<int>
accessed	2
accessory	2
accidentally	2
accidently	2
accounts	2
actions	2
activate	2
activated	2
activities	2
acts	2
1-10 of 872 rows	Previous 1 2 3 4 5 6 ... 88 Next

We see some different words now, but we still find these words usefull and very informative.

We now make variable only containing the words in descending order of how much they are used.

Hide

```
topwords <- reviews_tidy %>%
  count(word, sort = TRUE)

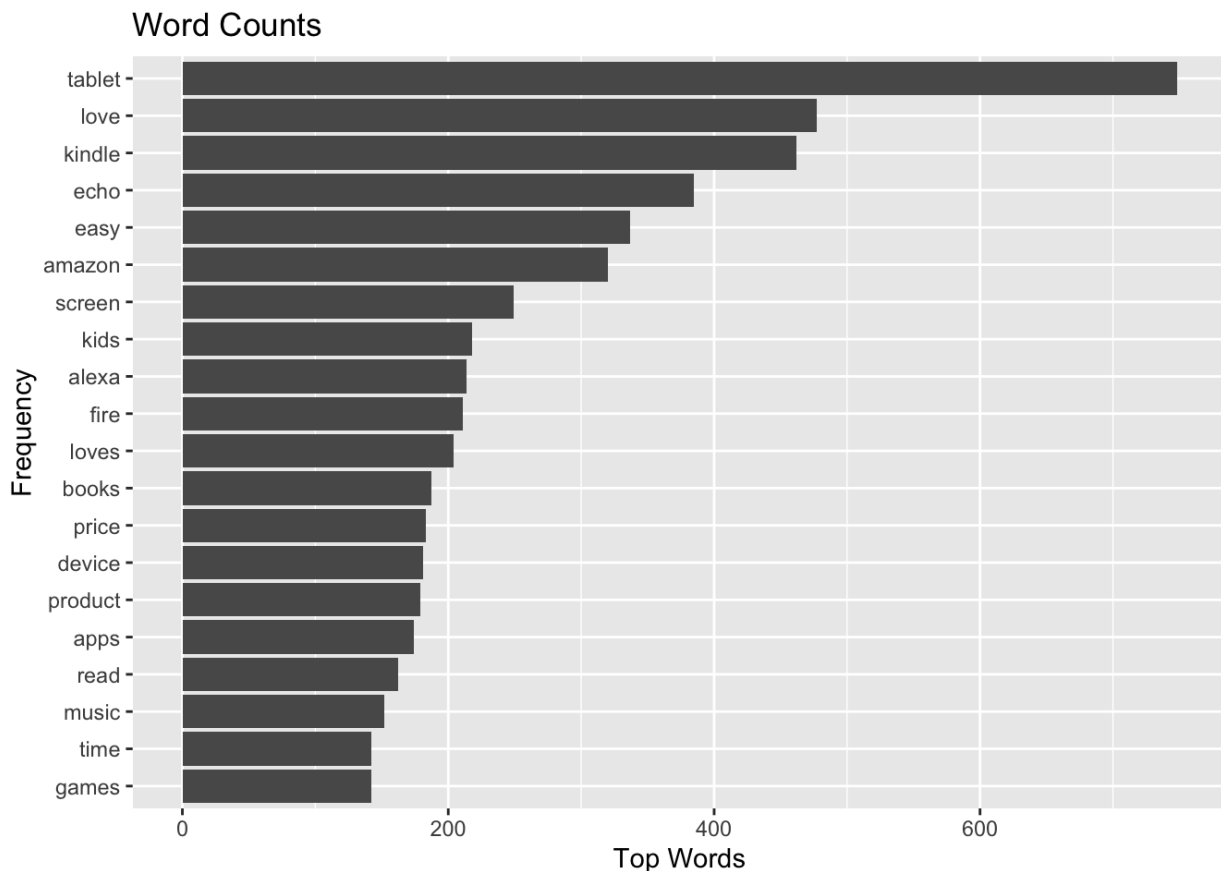
head(topwords)
```

<b>word</b>	<b>n</b>
<chr>	<int>
tablet	748
love	477
kindle	462
echo	385
easy	337
amazon	320
6 rows	

We use this new *topwords* variable to plot the topwords.

Hide

```
topwords %>%
  top_n(20, n) %>%
  ggplot(aes(x = word %>% fct_reorder(n), y = n)) + #fct_reorder factor the
counts + it reorder the counts in a descending format.
  geom_col() +
  coord_flip() +
  labs(title = "Word Counts",
        x = "Frequency",
        y = "Top Words")
```



We see that the word *tablet* is way more than the rest. If we have a look at the other words, we see that overall the words are mainly the name of some electronics or the words are mainly about electronics.

For later use, we save this *reviews\_tidy* with a different name

Hide

```
reviews_tidy_ML = reviews_tidy
```

## 2.2. Simple vectorization (Tf-idf)

In this section we add the *tf*, term of frequency, *idf*, inverse document frequency and *tf-idf*, term frequency–inverse document frequency to the tidy data.

Hide

```
reviews_tidy <- reviews_tidy %>%
  count(ID, word) %>%
  bind_tf_idf(ID, word, n) # I use this function to add the "tf", "idf" & "t
f_idf" values.
```

We have a look at the dwords with the highest *tf\_idf* values

[Hide](#)

```
reviews_tidy %>%
  arrange(desc(tf_idf))
```

ID	word	n	tf	idf	tf_idf
<int>	<chr>	<int>	<dbl>	<dbl>	<dbl>
102	ào	2	1.000000000	7.510431	7.510430556
238	âöre	2	1.000000000	7.510431	7.510430556
1167	response	2	1.000000000	7.510431	7.510430556
1250	bot	2	1.000000000	7.510431	7.510430556
1334	charges	2	1.000000000	7.510431	7.510430556
1412	mines	2	1.000000000	7.510431	7.510430556
2150	reliable	2	1.000000000	7.510431	7.510430556
2534	concept	2	1.000000000	7.510431	7.510430556
2881	launcher	2	1.000000000	7.510431	7.510430556
3094	lg	2	1.000000000	7.510431	7.510430556
1-10 of 7,828 rows			Previous	1	2
				3	4
				5	6
				...	783
				Next	

## 2.3. Topic modelling / Clustering (LSA)

Now we will perform a LSA, which is stable when attempting to do dimensionality reduction as preprocessing for supervised ML workflows, or for visualization.

We have loaded the quanteda package, which is for corpus-token based text analysis.

First we have to create a document-feature-matrix

[Hide](#)

```
reviews_dfm <- reviews_tidy %>%
  count(ID, word) %>%
  cast_dfm(document = ID, term = word, value = n)

reviews_dfm
```

```
## Document-feature matrix of: 1,510 documents, 1,827 features (99.7% sparse).
```

Now we get ready for the LSA, by choosing the number of dimensions.

[Hide](#)

```
reviews_lsa <- reviews_dfm %>%
  textmodel_lsa(nd = 5)
```

Let's have a look

Hide

```
reviews_lsa %>%
  glimpse()
```

```
## List of 5
## $ sk : num [1:5] 22.7 16.5 15 13.5 13
## $ docs : num [1:1510, 1:5] 0.0119 0.0191 0.0109 0.0199 0.015
...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:1510] "3" "4" "6" "7" ...
## .. ..$ : NULL
## $ features : num [1:1827, 1:5] 0.012388 0.004371 0.016513 0.000999
0.211243 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:1827] "dark" "didn't" "happy" "im" ...
## .. ..$ : NULL
## $ matrix_low_rank: num [1:1510, 1:1827] 0.0209 0.0266 0.0103 0.0177 0.02
58 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:1510] "3" "4" "6" "7" ...
## .. ..$ : chr [1:1827] "dark" "didn't" "happy" "im" ...
## $ data :Formal class 'dgCMatrix' [package "Matrix"] with 6 slots
## .. ..@ i : int [1:7828] 0 41 48 49 50 541 1395 0 959 1025 ...
## .. ..@ p : int [1:1828] 0 7 10 31 33 187 191 214 216 296 ...
## .. ..@ Dim : int [1:2] 1510 1827
## .. ..@ Dimnames:List of 2
## .. ..@ x : num [1:7828] 1 1 1 1 1 1 1 1 1 1 ...
## .. ..@ factors : list()
## - attr(*, "class")= chr "textmodel_lsa"
```

Now we take the *LSA documents* and put them into a tibble, and adding the id's as row names.

Hide

```
reviews_lsa_loading <- reviews_lsa$docs %>%
  as.data.frame() %>%
  rownames_to_column(var = "ID") %>%
  as_tibble()

reviews_lsa_loading %>%
  head()
```

ID	V1	V2	V3	V4	V5
<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
3	0.01185456	0.0037554460	0.01558929	0.04261870	0.01507959
4	0.01913180	0.0023359416	0.01422975	0.04757917	0.02522116
6	0.01092350	0.0000638529	0.01024060	0.01088557	0.01348919
7	0.01985597	0.0063195407	0.01095280	0.02735569	0.01406400
8	0.01501553	0.0028959729	0.01667973	0.04820387	0.02481269
11	0.04093220	0.0084140442	0.02277632	0.03733989	0.05381831

6 rows

We can nicely visualize it using *UMAP* dimensionality reduction for optimizing the visualization of the feature space.

Hide

```
reviews_lsa_umap = umap(reviews_lsa_loading %>%
  column_to_rownames("ID"),
  n_neighbors = 15,
  metric = "cosine",
  min_dist = 0.01,
  scale = TRUE,
  verbose = TRUE,
  n_threads = 8)
```

Then we put the result into a tibble for a nice layout.


Hide

```
reviews_lsa_umap = reviews_lsa_umap %>%
  as.data.frame()
```

Now we can plot the result

Hide

```
reviews_lsa_umap %>%
  ggplot(aes(x = V1, y = V2)) +
  geom_point(alpha = 0.5)
```



We see that the some of the points tend to cluster together.

We now try to find clusters by using the *hdbscan*.


Hide

```
reviews_lsa_hdbscan <- reviews_lsa_umap %>%
  as.matrix() %>%
  hdbscan(minPts = 50)
```

Let plot again, but now we color by clusters.

Hide

```
reviews_lsa_umap %>%
  bind_cols(cluster = reviews_lsa_hdbscan$cluster %>% as.factor(),
            prob = reviews_lsa_hdbscan$membership_prob) %>%
  ggplot(aes(x = V1, y = V2, col = cluster)) +
  geom_point(aes(alpha = prob), shape = 21)
```



## 2.4. EDA / simple frequency-based analysis

Now we look closer into rating, we want to see if some words tend to show more in one rating than the others.

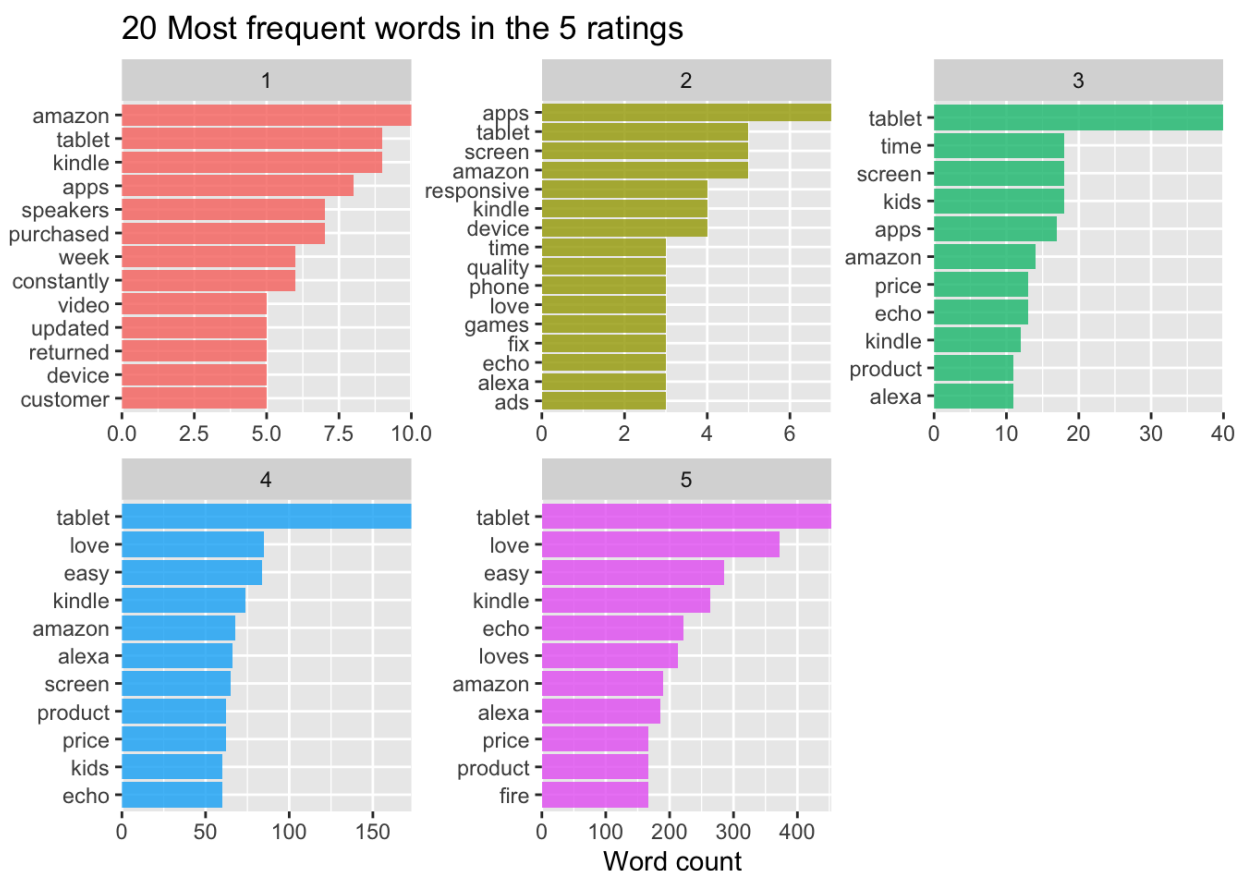
Hide

```
reviews_tidy3 = reviews_tidy %>%
  left_join(reviews, by = "ID")
```

Now we plot the 20 most used words in each rating.

Hide

```
reviews_tidy3 %>%
  count(rating, word, sort = TRUE) %>%
  group_by(rating) %>%
  top_n(10) %>%
  ungroup() %>%
  ggplot(aes(reorder_within(word, n, rating), n, fill = rating )) +
  geom_col(alpha = 0.8, show.legend = FALSE) +
  scale_x_reordered() +
  coord_flip() +
  facet_wrap(~rating, scales = "free") +
  scale_y_continuous(expand = c(0, 0)) +
  labs(x = NULL, y = "Word count", title = "20 Most frequent words in the 5 ratings")
```



we see that

- **Rating 1:** contains words like *constantly*, *returned*, *purchased* which could have a negative meaning and then there are a lot of words about the *electronics*.
- **Rating 2:** contains words like *responsive*, *time*, *fix* which also could have a negative meaning, and here we also see words describing the *electronics*
- **Rating 3:** contains words like the two other, but here the amount each words are used are almost the same for them all, except *tablet*. Therefore this looks like the ones that don't take this review seriously.
- **Rating 4:** contained words like *love*, *easy*, *price* and *product* this would usually have positive meaning in a text together with the description of the *electronic*.
- **Rating 5:** contain words like *love(s)*, *easy*, *price*, which are way more positive words than the first two ratings. We also see that the amount each words are used are way higher here.

We now want to see why the amount each words are used differ that much through the ratings.

Hide

```
reviews %>%
  group_by(rating) %>%
  count()
```

rating <fctr>	n <int>
1	63
2	54
3	197
4	1208
5	3478
5 rows	

We see that most of the reviewers gave the product a rating at 4 or 5, this could explain the differens in the amount of the word.

## 2.5. Topic modelling / Clustering (LDA)

### Preparing the Data

For this application, we have to leave the tidy data, since the *topicmodels* package requires a document-term matrix as input. We can easily produce it using the `cast_dtm()` function of *tidytext*. This matrix has to be term-frequency weighted, we do so using the *weightTf* function of the *tm* package for the weighting argument

Hide

```
reviews_dtm = reviews_tidy %>%
  count(ID, word) %>% #add a count of words in each id
  cast_dtm(document = ID, #Column containing document IDs as string or symbol
1
          term = word, #Column containing terms as string or symbol
          value = n, #Column containing values as string or symbol
          weighting = tm::weightTf) #The weighting function for the DTM/TDM

reviews_dtm
```

```
## <<DocumentTermMatrix (documents: 1510, terms: 1827)>>
## Non-/sparse entries: 7828/2750942
## Sparsity           : 100%
## Maximal term length: 18
## Weighting          : term frequency (tf)
```

We see that the matrix is 100% sparse with 1510 documents and 1827 terms, which is an artefact of text data generally. Now we try to see if we could reduce that somewhat by deleting less often used terms.

Hide



```
reviews_dtm %>% removeSparseTerms(sparse = .99) #sparse is maximal allowed sparsity.
```

```
## <<DocumentTermMatrix (documents: 1510, terms: 74)>>
## Non-/sparse entries: 3597/108143
## Sparsity           : 97%
## Maximal term length: 9
## Weighting           : term frequency (tf)
```

With max allowed sparsity at 99% we only have 74 terms out of 1827, too little

[Hide](#)

```
reviews_dtm %>% removeSparseTerms(sparse = .999) #sparse is maximal allowed sparsity.
```

```
## <<DocumentTermMatrix (documents: 1510, terms: 806)>>
## Non-/sparse entries: 6807/1210253
## Sparsity           : 99%
## Maximal term length: 13
## Weighting           : term frequency (tf)
```

With max allowed sparsity at 99.9% we only have 806 terms out of 1827, we might have to accept a high level of sparsity in order to still have a meaningful number of unique words.

Now we can perform an LDA, using the more accurate Gibbs sampling as method.

[Hide](#)

```
reviews_lda <- reviews_dtm %>%
  LDA(k = 5,
      method = "Gibbs",
      control = list(seed = 9))
```

## $\beta$ Word-Topic Association

$\beta$  is an output of the LDA model, indicating the propability that a word occurs in a certain topic. Therefore, looking at the top probability words of a topic often gives a good intuition regarding its properties.

[Hide](#)

```
# LDA output is defined for tidy(), so we can easily extract it
lda_beta <- reviews_lda %>%
  tidy(matrix = "beta") %>%
  group_by(topic) %>%
  arrange(topic, desc(beta)) %>% #the higher beta, the more a word is "used"
in a topic
  slice(1:15) %>%
  ungroup()

lda_beta %>% head()
```

topic	term
<int>	<chr>

beta
<dbl>

topic	term	beta
<int>	<chr>	<dbl>
1	easy	0.08183222
1	screen	0.04752102
1	kids	0.04694916
1	apps	0.03150912
1	device	0.02922171
1	set	0.02407503

6 rows


We see that topic one have words like *easy*, *screen* and *kids* with high  $\beta$  values.

Now we plot the top 10 words in each topic.

Hide

```
lda_beta %>%
  mutate(term = reorder_within(term, beta, topic)) %>% #reordering the term
  by beta in each topic
  group_by(topic, term) %>% #first group by topic and then by term
  arrange(desc(beta)) %>% #rearrange data in descending order by beta
  ungroup() %>%
  ggplot(aes(term, beta, fill = as.factor(topic))) + #aes for the plot
  geom_col(show.legend = FALSE) + #Don't show legend
  coord_flip() + #flip the graph, so the bars will be horizontal
  scale_x_reordered() +
  labs(title = "Top 10 terms in each LDA topic",
       x = NULL,
       y = expression(beta)) +
  facet_wrap(~ topic, nrow = 1, scales = "free")
```

Top 10 terms in each LDA topic



- **Topic 1:** We see that topic one is about quality for *HIFI* stuff.
- **Topic 2:** We see that topic two is mainly about reading. The word *Kindle* are top 1, but words like *books*, *read(ing)*, *light*, are also in here.
- **Topic 3:** We see that topic three is re mess of different types of words.
- **Topic 4:** We see that topic four is about *HIFI*, since words like *echo*, *alexa* and *music* are highly rated here.
- **Topic 5:** We see that topic five is mainly about tablets, since the words like *tablet*, *fire* and *ipad* are in this topic. together with some other tablet related words.

Overall we can say that the topics are mainly split into categories and not rating like we did earlier.

## 2.6. Embedding-model based vectorization (GloVe)

First we create a corpus

```
reviews_corpus <- reviews %>%
  corpus(docid_field = "ID",
        text_field = "text")
```

Hide

### Select features

Now we tokenize the corpus, we remove stopwords, punctuations, number and so on.

Hide

```
reviews_toks <- tokens(reviews_corpus, what = "word",
                      remove_numbers = TRUE, #Remove numbers
                      remove_punct = TRUE, #remove dots
                      remove_symbols = TRUE, #remove symbols
                      remove_separators = TRUE, #remove seperators
                      remove_url = TRUE, #remove url's
                      ngrams = 1) %>% #allowing unigrams
tokens_tolower() %>% #only lower case letters
tokens_remove(pattern = stopwords()) %>% #remove stopwords
tokens_remove(pattern = own_stopwords$word) %>% #remove own words
tokens_remove(pattern = c("[^[:alnum:]]")) #remove all non alpha numeric

reviews_toks %>% head(2)
```

```
## tokens from 2 documents.
## 1 :
## [1] "thought"      "big"          "small"        "paper"        "turn"
## [6] "just"         "like"         "palm"         "think"        "small"
## [11] "read"         "comfortable" "regular"      "kindle"       "definitely"
## [16] "recommend"    "paperwhite"   "instead"
##
## 2 :
## [1] "kindle"       "light"        "easy"         "use"          "especially"
## [6] "beach"
```

and then we keep the names of the features that occur five times or more.

Hide

```
feats <- dfm(reviews_toks, verbose = TRUE) %>%
  dfm_trim(min_termfreq = 5) %>%
  featnames()
```

## Constructing the feature co-occurence matrix

Hide

```
reviews_fcm <- fcm(reviews_toks,
                  context = "window",
                  count = "weighted",
                  weights = 1 / (1:5),
                  tri = TRUE)
```

## Fit word embedding model

Now we fit the *GloVe* model using *text2vec*

*GloVe* is an unsupervised learning algorithm for obtaining vector representations for words. Training is performed on aggregated global word-word co-occurrence statistics from a corpus, and the resulting representations showcase interesting linear substructures of the word vector space.

Hide

```
glove <- GlobalVectors$new(word_vectors_size = 50, #number of words in each vector

                           vocabulary = featnames(reviews_fcm),
                           x_max = 10)

reviews_main <- fit_transform(reviews_fcm,
                              glove,
                              n_iter = 20) #number of iterations
```

```
## INFO [2019-10-27 09:32:00] 2019-10-27 09:32:00 - epoch 1, expected cost 0.0806
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 2, expected cost 0.0479
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 3, expected cost 0.0389
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 4, expected cost 0.0336
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 5, expected cost 0.0300
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 6, expected cost 0.0273
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 7, expected cost 0.0251
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 8, expected cost 0.0234
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 9, expected cost 0.0220
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 10, expected cost 0.0207
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 11, expected cost 0.0197
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 12, expected cost 0.0188
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 13, expected cost 0.0180
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 14, expected cost 0.0173
## INFO [2019-10-27 09:32:01] 2019-10-27 09:32:01 - epoch 15, expected cost 0.0167
## INFO [2019-10-27 09:32:02] 2019-10-27 09:32:01 - epoch 16, expected cost 0.0161
## INFO [2019-10-27 09:32:02] 2019-10-27 09:32:02 - epoch 17, expected cost 0.0156
## INFO [2019-10-27 09:32:02] 2019-10-27 09:32:02 - epoch 18, expected cost 0.0152
## INFO [2019-10-27 09:32:02] 2019-10-27 09:32:02 - epoch 19, expected cost 0.0148
## INFO [2019-10-27 09:32:02] 2019-10-27 09:32:02 - epoch 20, expected cost 0.0144
```

## Averaging learned word vectors

The two vectors are main and context. According to the *Glove* paper, averaging the two word vectors results in more accurate representation.

[Hide](#)

```
reviews_context <- glove$components
dim(reviews_context)
```

```
## [1] 50 5613
```

We see that the dimension of this model is 50 x 5614.

[Hide](#)

```
reviews_vectors <- as.dfm(reviews_main + t(reviews_context))
```

## Examining term representations

Now we can find the six closest word vectors for *kindle*.

[Hide](#)

```
new <- reviews_vectors["kindle", ]

# calculate the similarity
cos_sim <- textstat_simil(reviews_vectors, new,
                          margin = "documents", method = "cosine")

head(sort(cos_sim[, 1], decreasing = TRUE), 6)
```

```
##   kindle   fire  second    new   much   first
## 1.0000000 0.8208233 0.6757720 0.6563131 0.6323784 0.6260871
```

We see the words *fire*, *new*, *since*, *one* and *second* are the closest words to the word *kindle*

The six closest word vectors for *tablet*.

[Hide](#)

```
new <- reviews_vectors["tablet", ]

# calculate the similarity
cos_sim <- textstat_simil(reviews_vectors, new,
                          margin = "documents", method = "cosine")

head(sort(cos_sim[, 1], decreasing = TRUE), 6)
```

```
##   tablet   great   good   price   fire  perfect
## 1.0000000 0.7494200 0.7299844 0.7288418 0.6668515 0.6633848
```

We see that the words like *good*, *great*, *price* and *kids* are the closest to the word *tablet*.

As seen earlier the data contains mostly of high rated reviews. Therefore it makes sense that the word *tablet* follows by positive words like *great* and *good*.

Now we make an average-vector-representations for the words

[Hide](#)

```
reviews_main_tibble = reviews_main %>%
  as.data.frame() %>% #making a dataframe
  rownames_to_column(var = "word") %>% #Row names
  as_tibble() #for nice layout

reviews_main_tibble %>% head()
```

word <chr>	V1 <dbl>	V2 <dbl>	V3 <dbl>	V4 <dbl>	V5 <dbl>	V6 <dbl>
thought	-0.2731706	-0.03957523	0.01527472	0.12807940	-0.08243837	-0.1331
big	0.1528246	-0.25964165	0.21220364	0.51244879	-0.30240378	-0.2765
small	-0.3029513	0.25653461	0.27191514	0.70423305	0.04605259	0.1813
paper	0.5156913	0.22531784	-0.23414992	0.18228483	0.01700168	-0.4546
turn	-0.5437862	-0.09267105	0.82648665	-0.01936661	0.22087295	0.3563
just	-0.4766783	0.01564204	-0.27068293	0.25873142	0.14141901	0.0895

6 rows | 1-7 of 51 columns

Now we make a average-vector-representations for the reviews.

Hide

```
reviews_tidy2 = reviews_toks %>%
  dfm() %>%
  tidy()

reviews_vector2 = reviews_tidy2 %>%
  inner_join(reviews_main_tibble, by = c("term" = "word")) #joining the data
  by word/term

head(reviews_vector2)
```

docum... <chr>	term <chr>	co... <dbl>	V1 <dbl>	V2 <dbl>	V3 <dbl>	V4 <dbl>	V5 <dbl>
1	thought	1	-0.2731706	-0.03957523	0.01527472	0.1280794	-0.08243
65	thought	1	-0.2731706	-0.03957523	0.01527472	0.1280794	-0.08243
70	thought	1	-0.2731706	-0.03957523	0.01527472	0.1280794	-0.08243
156	thought	1	-0.2731706	-0.03957523	0.01527472	0.1280794	-0.08243
158	thought	1	-0.2731706	-0.03957523	0.01527472	0.1280794	-0.08243
182	thought	1	-0.2731706	-0.03957523	0.01527472	0.1280794	-0.08243

6 rows | 1-8 of 53 columns

Hide

```
reviews_vector2 = reviews_vector2 %>%
  select(-term, -count) %>% #remove term and count
  group_by(document) %>% #group the data by document
  summarise_all(mean) %>% #finding the mean of the columns for each document
  ungroup()
```


Let's have a look at the data

Hide

```
reviews_vector2_umap <- umap(reviews_vector2 %>% column_to_rownames("document"),
                             n_neighbors = 15,
                             metric = "cosine",
                             min_dist = 0.01,
                             scale = TRUE,
                             verbose = TRUE) %>%
  as.data.frame()
```

Hide

```
reviews_vector2_umap %>%
  ggplot(aes(x = V1, y = V2)) +
  geom_point(shape = 21, alpha = 0.25)
```



It does not look like the data have any clusters right now. If we look closely we could argue about some few clusters at the plot, which might be *tablet* and/or *kindle* stuff.

### 3. Supervised / Unsupervised ML



## 3.1. Unsupervised ML

We do *kmeans* for the *tf\_idf* values, and want 5 clusters since we have 5 different ratings in the data.

[Hide](#)

```
kmeans_tfidf = kmeans(reviews_tidy[,4:6], centers = 5, nstart = 20)
```

[Hide](#)

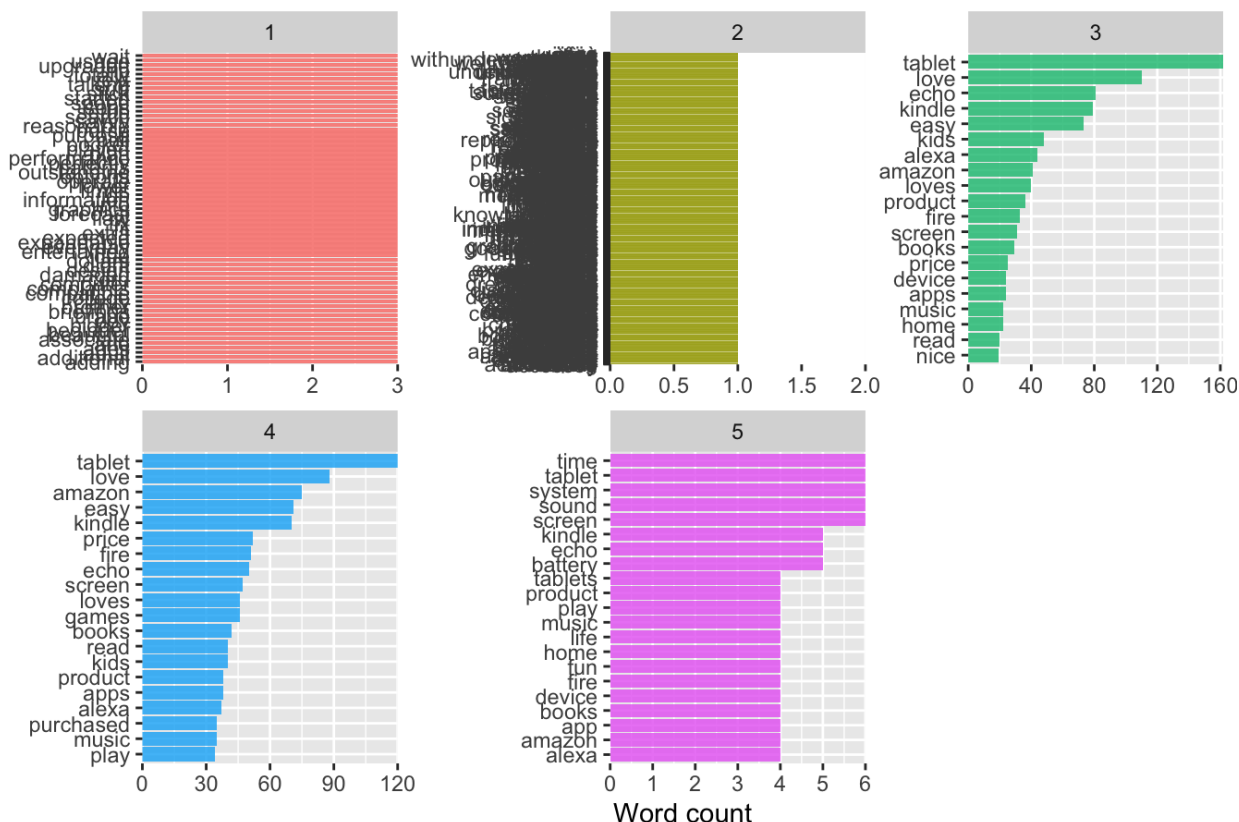
```
reviews_tidy = reviews_tidy %>%  
  bind_cols(cluster = kmeans_tfidf$cluster)
```

Now we plot the *reviews\_tidy*, and color by clusters.

[Hide](#)

```
reviews_tidy %>%  
  mutate(cluster = as.factor(cluster)) %>%  
  count(cluster, word, sort = TRUE) %>%  
  group_by(cluster) %>%  
  top_n(20,n) %>%  
  ungroup() %>%  
  ggplot(aes(reorder_within(word, n, cluster), n,  
    fill = cluster  
  )) +  
  geom_col(alpha = 0.8, show.legend = FALSE) +  
  scale_x_reordered() +  
  coord_flip() +  
  facet_wrap(~cluster, scales = "free") +  
  scale_y_continuous(expand = c(0, 0)) +  
  labs(x = NULL, y = "Word count",  
    title = "Most frequent words in the 5 clusters")
```

### Most frequent words in the 5 clusters



Unfortunately the plots are messy. since some of the clusters contain words represented the same amount of times. Therefore we can only talk/discuss plot 1, 2 and 5.

The plots contain almost the same words, it is hard to find differences in these three plots.

## 3.2. Supervised ML

### 3.2.1. Method 1

Now we want to get ready for modeling. We split the data into training and testing sets, using *rsample*. We use the original data, reviews, because the reviews\$text are our individual observations.

Hide

```
reviews_split <- reviews %>%
  select(ID2) %>%
  initial_split() #by default it splits 3/4

train_data <- training(reviews_split) #train data

test_data <- testing(reviews_split) #test data
```

Now we transform our training data from a tidy structure to a sparse matrix, which we will use for the machine learning algorithm.

Hide

```
sparse_words = reviews_tidy_ML %>%  
  count(ID2, word) %>%  
  inner_join(train_data) %>% #keep rows from Twwets_tidy where there are mat  
ching values in y, and all columns from x and y.  
  cast_sparse(ID2, word, n) #Create a sparse matrix from row names, column n  
ames, and values in a table
```

[Hide](#)

```
dim(sparse_words)
```

```
## [1] 1881 1754
```

We see we get a sparse matrix with 1885 rows and 1764 columns. This means we have 1885 training observations and 1764 features at this point.

Now we build a dataframe with response variable to associate each of the *rownames()* of the sparse matrix. We select the *rating*, since it will be our predictor later.

[Hide](#)

```
word_rownames = as.integer(rownames(sparse_words))  
reviews_rating = reviews %>% select(ID2, rating)
```

[Hide](#)

```
reviews_joined = data_frame(ID2 = word_rownames) %>%  
  left_join(reviews_rating, by = "ID2" )
```

Now we want to train our classification model. We use *glmnet* to fit a logistic regression model with LASSO regularization. It's a great fit for text classification because the variable selection that LASSO regularization performs can tell you which words are important for our prediction problem.

We make five models, since we have five types of ratings.

[Hide](#)

```
#Model for rate 1
model_1 = cv.glmnet(sparse_words, # x variable = matrix .
                    reviews_joined$rating == 1, #response variable
                    family = "binomial",
                    keep = TRUE)

#Model for rate 2
model_2 <- cv.glmnet(sparse_words, #x variable = matrix .
                    reviews_joined$rating == 2, #response variable
                    family = "binomial",
                    keep = TRUE)

#Model for rate 3
model_3 <- cv.glmnet(sparse_words, #x variable = matrix .
                    reviews_joined$rating == 3, #response variable
                    family = "binomial",
                    keep = TRUE)

#Model for rate 4
model_4 <- cv.glmnet(sparse_words, #x variable = matrix .
                    reviews_joined$rating == 4, #response variable
                    family = "binomial",
                    keep = TRUE)

#Model for rate 5
model_5 <- cv.glmnet(sparse_words, #x variable = matrix .
                    reviews_joined$rating == 5, #response variable
                    family = "binomial",
                    keep = TRUE)
```

Now we use *broom* to check the five models coefficients. We also check for which coefficients are largest in size, in each direction

[Hide](#)

```

coef_1 <- model_1$glmnet.fit %>%
  tidy() %>%
  filter(lambda == model_1$lambda.1se) %>% #lambda with error within 1 standard error
  group_by(estimate > 0) %>%
  top_n(10, abs(estimate)) %>%
  ungroup() %>%
  ggplot(aes(fct_reorder(term, estimate), #plotting the top words in each direction
              estimate,
              fill = estimate > 0)) +
  geom_col(alpha = 0.8, show.legend = FALSE) +
  coord_flip() +
  scale_fill_brewer(palette = "Paired") +
  labs(title = "Rate 1",
        x = NULL)

coef_2 <- model_2$glmnet.fit %>%
  tidy() %>%
  filter(lambda == model_2$lambda.1se) %>% #lambda with error within 1 standard error
  group_by(estimate > 0) %>%
  top_n(10, abs(estimate)) %>%
  ungroup() %>%
  ggplot(aes(fct_reorder(term, estimate),
              estimate,
              fill = estimate > 0)) +
  geom_col(alpha = 0.8, show.legend = FALSE) +
  coord_flip() +
  scale_fill_brewer(palette = "Paired") +
  labs(title = "Rate 2",
        x = NULL)

coef_3 <- model_3$glmnet.fit %>%
  tidy() %>%
  filter(lambda == model_3$lambda.1se) %>% #lambda with error within 1 standard error
  group_by(estimate > 0) %>%
  top_n(10, abs(estimate)) %>%
  ungroup() %>%
  ggplot(aes(fct_reorder(term, estimate),
              estimate,
              fill = estimate > 0)) +
  geom_col(alpha = 0.8, show.legend = FALSE) +
  coord_flip() + #Horizontal barplot
  scale_fill_brewer(palette = "Paired") +
  labs(title = "Rate 3",
        x = NULL)

coef_4 <- model_4$glmnet.fit %>%
  tidy() %>%
  filter(lambda == model_4$lambda.1se) %>% #lambda with error within 1 standard error
  group_by(estimate > 0) %>%

```

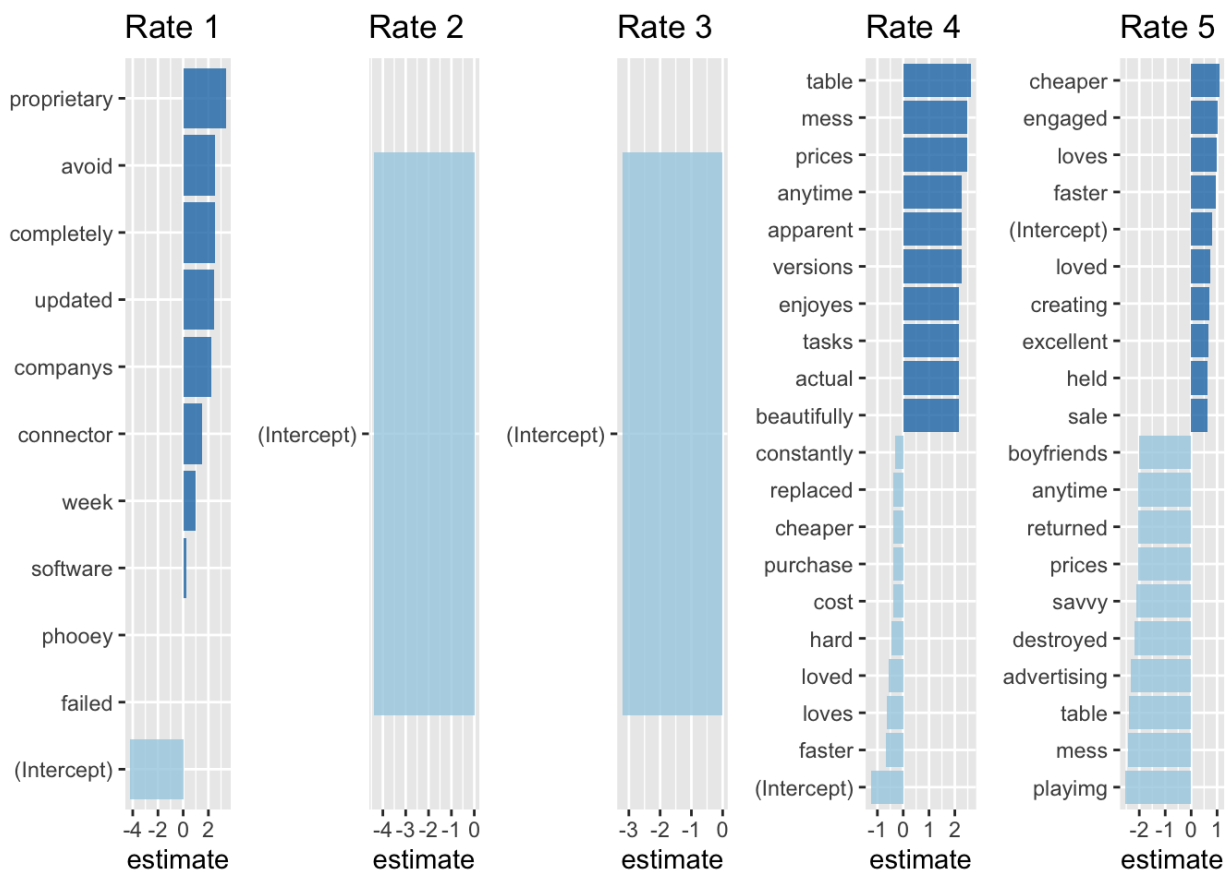
```

top_n(10, abs(estimate)) %>%
ungroup() %>%
ggplot(aes(fct_reorder(term, estimate),
             estimate,
             fill = estimate > 0)) +
geom_col(alpha = 0.8, show.legend = FALSE) +
coord_flip() + #Horizontal barplot
scale_fill_brewer(palette = "Paired") +
labs(title = "Rate 4",
      x = NULL)

coef_5 <- model_5$glmnet.fit %>%
  tidy() %>%
  filter(lambda == model_5$lambda.1se) %>% #lambda with error within 1 stand
ard error
group_by(estimate > 0) %>%
top_n(10, abs(estimate)) %>%
ungroup() %>%
ggplot(aes(fct_reorder(term, estimate),
             estimate,
             fill = estimate > 0)) +
geom_col(alpha = 0.8, show.legend = FALSE) +
coord_flip() + #Horizontal barplot
scale_fill_brewer(palette = "Paired") +
labs(title = "Rate 5",
      x = NULL)

grid.arrange(coef_1, coef_2, coef_3, coef_4, coef_5, nrow = 1)

```



We see that not all the plots are very good, but we can see the meaning of some words for some topics.

Now we try to classify the rating for the test set.

[Hide](#)

```
model <- cv.glmnet(sparse_words,
                  reviews_joined$rating,
                  family = "multinomial",
                  keep = TRUE)

coefs <- model$glmnet.fit %>%
  tidy() %>%
  filter(lambda == model$lambda.1se)

intercept <- coefs %>%
  filter(term == "(Intercept)") %>%
  pull(estimate)

classifications <- reviews_tidy_ML %>%
  inner_join(test_data) %>% #Predicting the test data
  inner_join(coefs, by = c("word" = "term")) %>%
  group_by(class, ID2) %>%
  summarize(score = sum(estimate)) %>%
  mutate(probability = plogis(0 + score))
```

Let's have a look.

[Hide](#)

classifications

class <chr>	ID2 <int>	score <dbl>	probability <dbl>
1	227	1.320891929	0.78933006
1	305	2.540961206	0.92696393
1	516	1.333006916	0.79133758
1	545	0.666503458	0.66071978
1	664	4.404218312	0.98792200
1	665	1.270480603	0.78082501
1	878	1.333006916	0.79133758
1	879	0.666503458	0.66071978
1	1782	0.666503458	0.66071978
1	1839	3.472099305	0.96988340
1-10 of 604 rows		Previous	1 2 3 4 5 6 ... 61 Next

We need to add the original rating, before we can say anything.

[Hide](#)

```
res <- classifications %>%
  group_by(ID2) %>%
  filter(probability == max(probability)) %>%
  left_join(reviews, by = "ID2")
```

Now we put the result in a table, where we can see if the classification did good or not.

[Hide](#)

```
result <- table(res$rating, res$class)

result
```

```
##
##      1      2      3      4      5
##  1      5      0      0      0      1
##  2      0      1      1      0      2
##  3      0      0     11      2      7
##  4      2      0      1     53     53
##  5      4      5     11     58    236
```

We see that the model is not that good, since it predicts  $\sim 50\%$  correct.

### 3.2.2. Method 2

Here we try do like we saw in M1.

We select our variables to the classification models.

[Hide](#)

```
reviews_tidy_ML <- reviews_tidy3 %>%
  select(tf, idf,
         tf_idf,
         helpful,
         rating,
         primaryCategories,
         year) %>%
  mutate(good_review = as.logical(as.integer(rating) >= 4, rating)) %>%
  select(-rating) # I set class == 0, so class 0 is TRUE and the rest FALSE:
```

### Splitting the data into train- and test set.

[Hide](#)

```
index <- createDataPartition(y = reviews_tidy_ML$good_review,
                             p = 0.75,
                             list = FALSE) # 75% to 25% split

training <- reviews_tidy_ML[index,]
test <- reviews_tidy_ML[-index,]
```

### Preprocessing using the recipes package.

[Hide](#)



```

reci <- recipe(good_review ~ ., data = training) %>%
  step_dummy(all_nominal(), -all_outcomes()) %>% # create dummy variables for the nominal variables
  step_center(all_numeric(), -all_outcomes()) %>% # Centers all numeric variables to mean = 0
  step_scale(all_numeric(), -all_outcomes()) %>% # this scales the numeric variables
  step_zv(all_predictors()) # Removed predictors with zero variance

reci %<>%
  prep(data = train)

```

## Setting the x and y values.

[Hide](#)

```

# Predictors
x_train <- bake(reci, new_data = training) %>% select(-good_review) # I remove class from the predictors.
y_train <- training %>% pull(good_review) %>% as.factor()
# test: split in y and x
x_test <- bake(reci, new_data = test) %>% select(-good_review)
y_test <- test %>% pull(good_review) %>% as.factor()

```

## Setting the workflow.

[Hide](#)

```

ctrl <- trainControl(method = "cv",
                     number = 4)
metric <- "Kappa" # I use Kappa because the class 0 (hate speech) is so low represented in the data.

```

## Fitting the model - logistic regression.

[Hide](#)

```

review_fit_log <- train(x = x_train,
                       y = y_train,
                       trControl = ctrl,
                       metric = metric,
                       method = "glm",
                       family = "binomial")

```

[Hide](#)

```
review_fit_log
```

```
## Generalized Linear Model
##
## 17694 samples
##      10 predictor
##      2 classes: 'FALSE', 'TRUE'
##
## No pre-processing
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 13270, 13271, 13270, 13271
## Resampling results:
##
##      Accuracy   Kappa
##      0.9206511   0
```

## Prediction and evaluating.

[Hide](#)

```
pred_log <- predict(review_fit_log, newdata = x_test)
```

[Hide](#)

```
print("Confusion matrix for logistic model")
```

```
## [1] "Confusion matrix for logistic model"
```

[Hide](#)

```
confusionMatrix(pred_log, y_test, positive = 'TRUE')
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE      0      0
##      TRUE     468 5429
##
##           Accuracy : 0.9206
##           95% CI : (0.9134, 0.9274)
##      No Information Rate : 0.9206
##      P-Value [Acc > NIR] : 0.5123
##
##           Kappa : 0
##
##  Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 1.0000
##           Specificity : 0.0000
##      Pos Pred Value : 0.9206
##      Neg Pred Value :      NaN
##           Prevalence : 0.9206
##      Detection Rate : 0.9206
##      Detection Prevalence : 1.0000
##      Balanced Accuracy : 0.5000
##
##           'Positive' Class : TRUE
##
```

## Fitting the model - decision tree

[Hide](#)

```
reviews_fit_dt <- train(x = x_train,
                        y = y_train,
                        trControl = ctrl,
                        metric = metric,
                        method = "rpart")
```

[Hide](#)

```
reviews_fit_dt
```

```
## CART
##
## 17694 samples
##      10 predictor
##      2 classes: 'FALSE', 'TRUE'
##
## No pre-processing
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 13270, 13271, 13271, 13270
## Resampling results across tuning parameters:
##
##      cp          Accuracy   Kappa
##      0.01495726  0.9434272  0.4251322
##      0.02207977  0.9299197  0.1796601
##      0.02250712  0.9299197  0.1796601
##
## Kappa was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.01495726.
```

## Predicting and evaluating

[Hide](#)

```
pred_dt <- predict(reviews_fit_dt, newdata = x_test)
```

[Hide](#)

```
print("Confusion matrix for decision tree")
```

```
## [1] "Confusion matrix for decision tree"
```

[Hide](#)

```
confusionMatrix(pred_dt, y_test, positive = 'TRUE')
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE   134    0
##      TRUE    334 5429
##
##              Accuracy : 0.9434
##              95% CI : (0.9372, 0.9491)
##      No Information Rate : 0.9206
##      P-Value [Acc > NIR] : 7.679e-12
##
##              Kappa : 0.4249
##
##      Mcnemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 1.0000
##              Specificity : 0.2863
##      Pos Pred Value : 0.9420
##      Neg Pred Value : 1.0000
##      Prevalence : 0.9206
##      Detection Rate : 0.9206
##      Detection Prevalence : 0.9773
##      Balanced Accuracy : 0.6432
##
##      'Positive' Class : TRUE
##
```

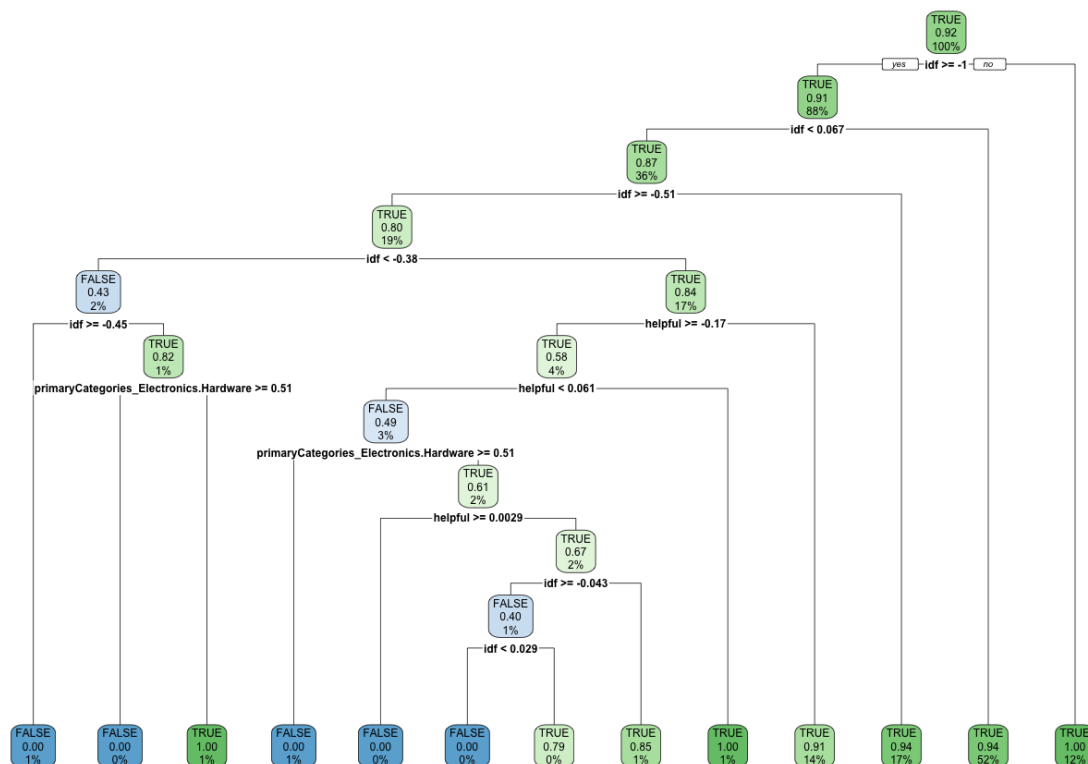
*We had some problems with the margins therefore we changed it.*

Hide

```
par(mar = c(1,1,1,1)) #The margins for the next plot are too large
```

Hide

```
reviews_fit_dt$finalModel %>%
  rpart.plot()
```



## Fitting the model - Random forest

Hide

```
reviews_fit_rf <- train(x = x_train,
                        y = y_train,
                        trControl = ctrl,
                        metric = metric,
                        method = "ranger",
                        importance = "impurity", # To define how to measure variable
                        importance (later more)
                        num.trees = 25
                        )
```

Hide

```
reviews_fit_rf
```

```
## Random Forest
##
## 17694 samples
##    10 predictor
##    2 classes: 'FALSE', 'TRUE'
##
## No pre-processing
## Resampling: Cross-Validated (4 fold)
## Summary of sample sizes: 13271, 13270, 13270, 13271
## Resampling results across tuning parameters:
##
##   mtry  splitrule  Accuracy  Kappa
##   2     gini       0.9255679  0.1072797
##   2     extratrees 0.9206511  0.0000000
##   6     gini       0.9540520  0.6316107
##   6     extratrees 0.9511698  0.5874154
##  10     gini       0.9499264  0.6216231
##  10     extratrees 0.9502091  0.6222944
##
## Tuning parameter 'min.node.size' was held constant at a value of 1
## Kappa was used to select the optimal model using the largest value.
## The final values used for the model were mtry = 6, splitrule = gini
##   and min.node.size = 1.
```

## Predicting and evaluating

[Hide](#)

```
pred_rf <- predict(reviews_fit_rf, newdata = x_test)
```

[Hide](#)

```
print("Confusion matrix for Random forest")
```

```
## [1] "Confusion matrix for Random forest"
```

[Hide](#)

```
confusionMatrix(pred_rf, y_test, positive = 'TRUE')
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction FALSE TRUE
##      FALSE    272    56
##      TRUE     196  5373
##
##           Accuracy : 0.9573
##           95% CI : (0.9518, 0.9623)
##      No Information Rate : 0.9206
##      P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6613
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.9897
##           Specificity : 0.5812
##      Pos Pred Value : 0.9648
##      Neg Pred Value : 0.8293
##           Prevalence : 0.9206
##      Detection Rate : 0.9111
##      Detection Prevalence : 0.9444
##      Balanced Accuracy : 0.7854
##
##           'Positive' Class : TRUE
##
```

## 4. Network Analysis

In this section we are going to construct a network based on the reviews data.

We did several constructions and tried interpret different attributes to the network, but for some reason we could not include any attributes. For example we wanted to try include *rating* and *primaryCategories* and use them for coloring or shaping the nodes

### Creating smaller sample of reviews.

We first tried to create a network including the whole dataset, but because of the size, we where not be able to run the network/graph (*the system crashed*).

Instead we resampled the review dataset to only 500 observations - so we instead got a datasat consiting of 500 reviews instead of 5.000.

Hide

```
reviews_sampled <- reviews[sample(nrow(reviews),500),]
```

Now we create a tidy data as we did earlier, just on the resampled set.

Hide



```
reviews_sampled_tidy <- reviews_sampled %>%
  select(ID, text) %>%
  unnest_tokens(output = word, input = text) %>%
  anti_join(stop_words %>% bind_rows(own_stopwords), by = "word") %>%
  mutate(word = trimws(gsub("[^\\s]*[0-9][^\\s]*", "", word, perl = T))) %>%
  filter(str_length(word) > 1) %>% # this filter out the words that are blank.
  mutate(word = word %>% str_remove_all("[^[:alnum:]]") ) %>% # alnum = Alphanumeric characters.
  filter(str_length(word) > 1) # filter out words with 1 character.
```

To make a network not with too many words, we only use the 50 highest counted words in the tidy set.

Hide

```
sampled_50 <- reviews_sampled_tidy %>%
  count(word, sort = TRUE) %>% # we count the words
  head(50) # we take the head of the 50 most frequent words.
```

Now we want the sampled tidy to only consist of the 50 most frequent words and here we use “left\_join” to this.

First we get a big dataset, but this is because we have all the words with “NA” values, which is the words that were not a part of the 50 most frequent words. Therefore we use *drop\_na()* to drop the rows with NA values.

Hide

```
reviews_sampled_tidy_joined <- reviews_sampled_tidy %>%
  left_join(sampled_50, by = "word") %>%
  drop_na()
```

Let's have a look.

Hide

```
glimpse(reviews_sampled_tidy_joined)
```

```
## Observations: 1,800
## Variables: 3
## $ ID      <int> 1825, 1825, 1825, 3594, 3594, 817, 817, 817, 817, 577, 577,...
## $ word    <chr> "tablet", "tablet", "games", "recommend", "kindle", "alexa"...
## $ n       <int> 127, 127, 39, 22, 89, 40, 98, 19, 98, 21, 46, 25, 42, 16, 6...
```

We see that the data now consists of 1933 observations and 3 variables.

Now we can create the nodelist from the sampled tidy dataset.

Hide

```
nodes_reviews <- reviews_sampled_tidy_joined %>%
  group_by(ID) %>%
  select(word, ID) %>%
  ungroup() %>%
  distinct(word) # distinct does so we only retain the unique words from the nodelist -> so we don't get duplicate words/nodes.
```

Then we create a dataframe that we are going to be use for *left\_joining* the edges (word.x and word.y)

Hide

```
nodes4join <- reviews_sampled_tidy_joined %>%  
  select(ID, word)
```

Hide

```
g_reviews <- reviews_sampled_tidy_joined %>%  
  left_join(nodes4join, by = c("ID" = "ID"))
```

Hide

```
g_reviews <- g_reviews[c("word.x", "word.y")]
```

Here we transform the dataframe into a graph using *igraph*.

We set *directed = FALSE*, because it is an undirected network.

Hide

```
g_reviews <- graph_from_data_frame(d = g_reviews,  
                                  directed = FALSE,  
                                  vertices = nodes_reviews)
```

We use *simplify* to remove loops and multiple edges.

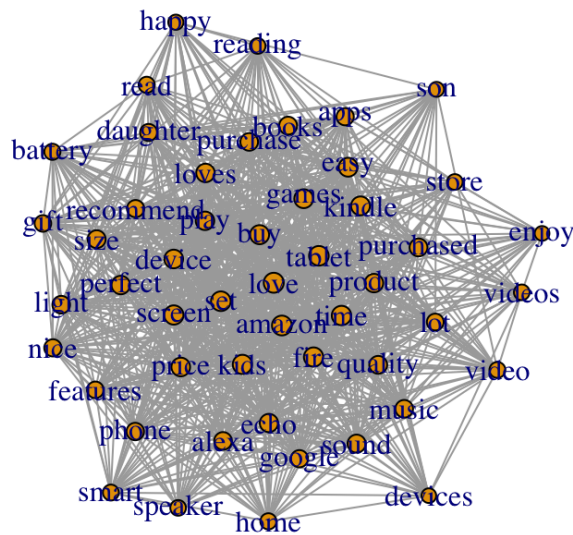
Hide

```
g_reviews <- simplify(g_reviews,  
                     remove_multiple = TRUE,  
                     remove_loops = TRUE,  
                     edge_attr_comb = igraph_opt("edge.attr.comb"))
```

Now we can plot the network.

Hide

```
plot(g_reviews,  
     vertex.size = 1 + sqrt(degree(g_reviews, mode = "all"))) # we assign the  
     size to the degree and add "1+sqrt", so we dont get too big nodes.
```



At the graph we see that the size based on the degree is not that different between the words, but it seems as *device*, *tablet* & *love* is connected by *ID* to many words. This is not that surprising, because these were also some of the topwords from earlier.

Lets have a look at each nodes degree.

Hide

```
degree(g_reviews) %>%  
  as.data.frame()
```

	<dbl>
tablet	47
games	43
recommend	35
kindle	43
alexa	37
love	49
features	32
video	29
screen	45
perfect	41

1-10 of 50 rows

Previous **1** 2 3 4 5 Next

We can see that fx. *tap* only has degree 25, but *love* has degree 49. This makes sense, that not that many reviews has *tap* in the text, but rather more has *love* inside (*probably because the many high rated reviews*).

## Alternative solution

We could have looked into bigrams and plot those as a network.

Hide

```
bigrams <- reviews %>%
  unnest_tokens(bigram, text, token = "ngrams", n = 2) # n is the number of
  words to consider
```

Hide

```
bigrams_separated <- bigrams %>%
  separate(bigram, c("word1", "word2"), sep = " ")
```

Hide

```
bigrams_filtered <- bigrams_separated %>%
  filter(!word1 %in% stop_words$word) %>%
  filter(!word2 %in% stop_words$word) %>%
  filter(!word1 %in% own_stopwords$word) %>%
  filter(!word2 %in% own_stopwords$word)
```

Hide

```
bigram_counts <- bigrams_filtered %>%
  count(word1, word2, sort = TRUE)
```

Hide


```
bigram_counts %>%
  head()
```

word1 <chr>	word2 <chr>	n <int>
kindle	fire	145
battery	life	100
amazon	fire	78
amazon	echo	76
amazon	prime	66
smart	home	66
6 rows		

Hide

Hide

degree(bigram\_graph)



- 1
- 2
- 3
- 4
- 5
- 6